Homework 2 — Linear Regression for Age Estimation (preview on subset)
Subset sizes — train: 2000 / 5000, test: 2000 / 2500
Pixels scaled to [0,1]; optimizer: mini-batch SGD (B=256).

Method	train half-MSE	test half-MSE
SGD (no reg)	109.851502	111.190394
SGD (L2, alpha=0.1)	109.851554	111.190323

(d) L2 model (alpha=0.1): Test RMSE (subset) = 14.9124 years

Weights — SGD (no reg)

Weights — SGD (L2, alpha=0.1)

Worst error #1
GT=89.00, Pred=39.35, |err|=49.65

Worst error #2
GT=84.50, Pred=39.47, |err|=45.03

Worst error #3
GT=80.00, Pred=39.36, |err|=40.64

Worst error #4
GT=80.00, Pred=39.46, |err|=40.54

Worst error #5
GT=80.00, Pred=39.47, |err|=40.53